

```
In [1]: ▶ #Import neccesary Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```


```
In [2]: ▶ #Load dataset and show top 5 rows
df = pd.read_csv("heart.csv")
df.head ()
```

Out[2]:

	SN	biking	smoking	heartdisease
0	1	30.801246	10.896608	11.769423
1	2	65.129215	2.219563	2.854081
2	3	1.959665	17.588331	17.177803
3	4	44.800196	2.802559	6.816647
4	5	69.428454	15.974505	4.062224

```
In [3]: ▶ #scatterplot heart disease vs. biking
y = df['heartdisease']
X = df['biking']
sns.scatterplot(x=X,y=y)
```


Out[3]: <AxesSubplot:xlabel='biking', ylabel='heartdisease'>



```
In [4]: ▶ #This scatter plot reveals that there is a highly negative
#correlation between individuals who bike and the risk of
#heart disease. This will reflect on the model by providing
#accurate measures of correlation with the subject matter.
```

```
In [5]: ▶ #scatterplot heart disease vs. smoking.  
y = df['heartdisease']  
X = df['smoking']  
sns.scatterplot(x=X,y=y)
```


Out[5]: <AxesSubplot:xlabel='smoking', ylabel='heartdisease'>



```
In [6]: ▶ #This scatter plot reveals that there is a weak positive  
#correlation between individuals who smoke and the risk  
#of heart disease. This will reflect on the model by prompting  
#to conduct more analysis between smoking and heart disease  
#to further collect more information on the complexity within  
#the correlation
```

```
In [7]: ▶ #scatterplot smoking vs biking  
y = df['smoking']  
X = df['biking']  
sns.scatterplot(x=X,y=y)
```

Out[7]: <AxesSubplot:xlabel='biking', ylabel='smoking'>



```
In [8]: ▶ #This scatter plot reveals that there is no correlation between  
#individuals who smoke and the individuals who bike. This will reflect  
#on the model by reporting that further analysis drawn between biking  
#and smoking will not be necessary for the remainder of the model.
```

```
In [9]: > #drop unnecessary columns
X = df.drop(['heartdisease', 'SN'],axis=1)
y = df['heartdisease']
```

```
In [10]: > #split data into training and test sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=15)
```

```
In [11]: > #the purpose of splitting is to avoid the issue of over and under fitting
#so we can train the model by using the training set and apply it to
#properly evaluate the performance of the model
```

```
In [12]: > #Training a Linear regression model
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train,y_train)
```

Out[12]: LinearRegression()

```
In [13]: > #display coefficients of the regression model
modelCoef = model.coef_
coeff_df = pd.DataFrame(modelCoef,X.columns,columns=['Coefficient'])
coeff_df
```

Out[13]:

	Coefficient
biking	-0.200524
smoking	0.176604

```
In [14]: > #Calculating the coefficient is a way to define to relationship between
#predictor variables and the response variables
```

```
In [15]: > #The metrics that are going to be computed are the mean absolute error,
#mean squared error, R^2 and the root mean square error
```

```
In [16]: > #generate test predictions using the model
test_predictions = model.predict(X_test)
```

```
In [17]: > #import metrics and calculate MAE, MSE, and RMSE
from sklearn.metrics import mean_absolute_error,mean_squared_error
MAE = mean_absolute_error(y_test,test_predictions)
MSE = mean_squared_error(y_test,test_predictions)
RMSE = np.sqrt(MSE)
```

```
In [18]: > MAE
```

Out[18]: 0.5059659645518587

```
In [19]: > MSE
```

Out[19]: 0.41549278914354126

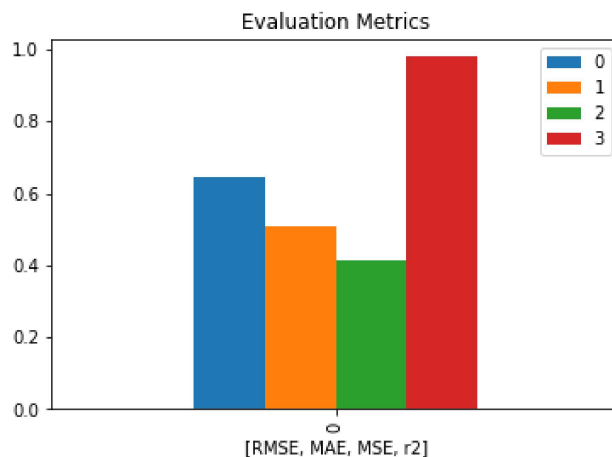
In [20]: ▶ RMSE

Out[20]: 0.6445873014135023

In [21]: ▶ *#calculate and display r-squared value*  
`from sklearn.metrics import r2_score`  
`r2=r2_score(y_test, test_predictions)`  
`r2`

Out[21]: 0.9792389107198726

In [22]: ▶ *#Create a bar chart to visualize RMSE, MSE, MAE, r2*  
`RMSE = pd.Series(RMSE)`  
`MSE = pd.Series(MSE)`  
`MAE = pd.Series(MAE)`  
`r2 = pd.Series(r2)`  
`metrics = pd.concat([RMSE, MAE, MSE, r2], axis = 1)`  
`metrics.plot.bar(xlabel=['RMSE', 'MAE', 'MSE', 'r2'])`  
`plt.title('Evaluation Metrics')`  
`plt.show()`



In [23]: ▶ *#The root square mean error is .644 which means the model can predict the data accurately because the measure falls between 0.2 and 0.6.*  
*#The r2 is .979 which means this model is reading a high correlation between the independent data and the dependent data..*  
*#The mean absolute error is .505 which means the data is medially reliable because the closer the number is to 0, the more reliable the model is.*  
*#The mean squared error is .415 meaning the model is reliable because the closer this measure is to 0, the less error the model has.*

In [24]: ▶ *#Create a DataFrame for new subjects and predict heart disease based on biking and smoking data*  
`subjects = [[3.99,13.00], [19.11,15.17], [44.43,19.66],`  
`[55.13,7.05],[34.22,15.05]]`  
`subjects = pd.DataFrame(subjects, columns = ['biking', 'smoking',])`  
`model.predict(subjects)`

Out[24]: array([16.51119814, 13.86249862, 9.57817036, 5.20558621, 10.81138186])